Eleanor collected and cleaned data from public sources, worked on the data simulation, worked on the Smith-Waterman algorithm, and did time/memory evaluation. Matthew developed our data-utils functions, worked on MinHash and integrated the different methods into a single pipeline. Gabe worked on the data simulation and programmed the k-mer index method. Alexey created our FM index based method. All members wrote their respective parts in the writeup and helped with editing the document.